

BD_24_Summative - Final Rev 2

July 20, 2018

1 Summative - Dieter Gebert

Final thoughts at the end

This exercise will go through a simulated streaming data workflow, as seen in the streaming data module. But it will be more complex, and closer to what a real world scenario might look like. You will need to - Investigate the incoming data - Create an appropriate database to store the incoming records - Write code to process records one by one as they arrive, including - Printing warnings when any reading goes above a predefined threshold - Storing the incoming data in a database - Write code to analyse the stored data in a scalable manner - Display relevant information in a dashboard

1.1 Create an appropriate database

Each incoming record will look something like the following:

```
{'Device_ID': 9,  
'Temp1': 33.01235436945101,  
'Temp2': 46.313589806396116,  
'Temp3': 16.506177184725505,  
'Temp_Ambient': 23.782493817278034}
```

Each device is assigned an integer ID. Every device has multiple sensors, and reports the readings from each sensor as a float. Each call to `gen_data.getReading()` returns a time (an integer here to make things easier) and a record that follows the same pattern as above.

Create a database to store the incoming data. Include a time field for the time that the data arrives. If you create the database with python, show the code here, otherwise include any bash or sql code you run. You may wish to come back and add additional fields to make later analysis easier.

1.2 Store and process the incoming data

As each record arrives (i.e. each loop of the for loop), you must - Store the record in the database you created above - Use either moving windows or exponential averaging to keep track each sensor value for each device. Print out the values at the end of the loop - Print a warning if any reported temperature exceeds 100 degrees for the first time for that device

```
In [1]: import pandas as pd  
import numpy as np
```

```

In [2]: import gen_data

        # Your code here for any initializations you may need

temp1_ave = 0
temp2_ave = 0
temp3_ave = 0
temp_am_ave = 0

window1 = []
window2 = []
window3 = []
windowA = []

window_t1_sd = 0
window_t1_av = 0
window_t2_sd = 0
window_t2_av = 0
window_t3_sd = 0
window_t3_av = 0
window_ta_sd = 0
window_ta_av = 0

df = pd.DataFrame(columns=['Time', 'Device_ID', 'Temp1', 'Temp1_ave', 'Temp1_sd',
                           'Temp2', 'Temp2_ave', 'Temp2_sd',
                           'Temp3', 'Temp3_ave', 'Temp3_sd',
                           'Temp_Ambient', 'Temp_Am_ave', 'Temp_Am_sd'])

for i in range(2000): # Hint: make this lower for testing original value: 20000

    # The simulated data arriving - don't change this
    arrival_time, record = gen_data.getReading()

    # Your code here
    temp1_ave = temp1_ave*0.75 + record['Temp1']*0.25
    temp2_ave = temp2_ave*0.75 + record['Temp2']*0.25
    temp3_ave = temp3_ave*0.75 + record['Temp3']*0.25
    temp_am_avg = temp_am_ave*0.75 + record['Temp_Ambient']*0.25

    # Window for temp1
    window1.append(temp1_ave) # Add the temp1_ave to our moving window
    if len(window1)>10: # Keep the window size from growing beyond 10:
        del(window1[0]) # If the window is >10 items, delete the oldest
    window_t1_sd = np.std(window1) # Calculate the standard deviation of the ten items in window
    window_t1_av = np.mean(window1) # Calculate the mean of the last ten readings

    # Window for temp2
    window2.append(temp2_ave) # Add the temp2_ave to our moving window

```

```

if len(window2)>10: # Keep the window size from growing beyond 10:
    del(window2[0]) # If the window is >10 items, delete the oldest
window_t2_sd = np.std(window2) # Calculate the standard deviation of the ten items i
window_t2_av = np.mean(window2) # Calculate the mean of the last ten readings

# Window for temp3
window3.append(temp3_ave) # Add the temp3_ave to our moving window
if len(window3)>10: # Keep the window size from growing beyond 10:
    del(window3[0]) # If the window is >10 items, delete the oldest
window_t3_sd = np.std(window3) # Calculate the standard deviation of the ten items i
window_t3_av = np.mean(window3) # Calculate the mean of the last ten readings

# Window for temp_ambient
windowA.append(temp_am_avg) # Add the temp_am_ave to our moving window
if len(windowA)>10: # Keep the window size from growing beyond 10:
    del(windowA[0]) # If the window is >10 items, delete the oldest
window_ta_sd = np.std(windowA) # Calculate the standard deviation of the ten items i
window_ta_av = np.mean(windowA) # Calculate the mean of the last ten readings

# Writing simulated data to df
df = pd.concat([df, pd.DataFrame([{'Time': arrival_time, 'Device_ID': record['Device
                                'Temp1':record['Temp1'], 'Temp1_ave':window_t1_av,
                                'Temp2':record['Temp2'], 'Temp2_ave':window_t2_av,
                                'Temp3':record['Temp3'], 'Temp3_ave':window_t3_av,
                                'Temp_Ambient':record['Temp_Ambient'], 'Temp_Am_av

# Going though the df to check for Temperatures that are out of bound
for g in range(4):
    if df.iloc[g][1] >= 100 :
        print('Warning! Deive ID:', df.iloc[g][0], 'Temp1 is above 100 řC. Reading: ', df
    elif df.iloc[g][2] >= 100 :
        print('Warning! Deive ID:', df.iloc[g][0], 'Temp2 is above 100 řC. Reading: ', df
    elif df.iloc[g][3] >= 100 :
        print('Warning! Deive ID:', df.iloc[g][0], 'Temp3 is above 100 řC. Reading: ', df
    elif df.iloc[g][4] >= 100 :
        print('Warning! Deive ID:', df.iloc[g][0], 'Temp_Ambient is above 100 řC. Reading

```

I did not like the arrival_time recorded so I used the index as my arrival time instead

```
In [3]: df.reset_index(level=0, inplace=True)
```

create a datetime from the index

- epoch conversion - the origion is fixed though
- re-arrange the df

```
In [4]: df['Date'] = pd.to_datetime(df['index'], unit='s', origin = '2018-07-22')
```

```
df = df[['Date', 'Device_ID', 'Temp1', 'Temp1_ave', 'Temp1_sd',
        'Temp2', 'Temp2_ave', 'Temp2_sd',
        'Temp3', 'Temp3_ave', 'Temp3_sd',
        'Temp_Ambient', 'Temp_Am_ave', 'Temp_Am_sd']]
```

```
# Show data frame
df.head(5)
```

```
Out[4]:
```

	Date	Device_ID	Temp1	Temp1_ave	Temp1_sd	Temp2	\
0	2018-07-22 00:00:00	4	26.809323	6.702331	0.000000	39.490255	
1	2018-07-22 00:00:01	7	47.118053	11.754296	5.051965	61.829858	
2	2018-07-22 00:00:02	6	42.906642	15.613316	6.840978	57.197306	
3	2018-07-22 00:00:03	4	41.089176	18.652690	7.925441	55.198093	
4	2018-07-22 00:00:04	8	58.213078	21.998428	9.748124	74.034385	

	Temp2_ave	Temp2_sd	Temp3	Temp3_ave	Temp3_sd	Temp_Ambient	\
0	9.872564	0.000000	13.404661	3.351165	0.000000	21.938834	
1	16.367226	6.494662	23.559027	5.877148	2.525983	22.103256	
2	21.393398	8.868215	21.453321	7.806658	3.420489	22.052069	
3	25.391006	10.340529	20.544588	9.326345	3.962720	21.742404	
4	29.622098	12.535941	29.106539	10.999214	4.874062	22.011406	

	Temp_Am_ave	Temp_Am_sd
0	5.484709	0.000000
1	5.505261	0.020553
2	5.507847	0.017175
3	5.489785	0.034639
4	5.492398	0.031420

Writing to database - csv file

```
In [5]: with open("sensor_data.csv","a") as f:
        df.to_csv(f,header=True,index=False)
```

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 14 columns):
Date                2000 non-null datetime64[ns]
Device_ID           2000 non-null object
Temp1               2000 non-null float64
Temp1_ave           2000 non-null float64
Temp1_sd            2000 non-null float64
Temp2               2000 non-null float64
Temp2_ave           2000 non-null float64
Temp2_sd            2000 non-null float64
Temp3               2000 non-null float64
```

```
Temp3_ave      2000 non-null float64
Temp3_sd       2000 non-null float64
Temp_Ambient   2000 non-null float64
Temp_Am_ave    2000 non-null float64
Temp_Am_sd     2000 non-null float64
dtypes: datetime64[ns](1), float64(12), object(1)
memory usage: 218.8+ KB
```

1.3 Analyzing the stored data

You now have a nice big database. Load it into spark for analysis.

You are told that during the time the data was being collected, devices 3 and 10 had malfunctioning sensors - their temperature3 readings are all 200+. Verify this. Since the engineers knew about the faulty sensors, no harm has been done, but seeing those false readings in the historical data makes you unhappy. You decide to go the extra mile and replace these readings with slightly more believable (but still false) data, to practise your new machine learning skills.

Using the other devices for training, build a model to predict temperature3 given readings from the other sensors. Use the model to replace the erroneous values with the predicted ones.

Do you think this is a reasonable step to take? Explain.

```
In [69]: import findspark
         findspark.init()
```

```
#from pyspark.sql.session import SparkSession
from pyspark.sql import SparkSession
```

```
spark = SparkSession.builder.master("local[*]").getOrCreate()
```

```
ERROR:py4j.java_gateway:An error occurred while trying to connect to the Java server (127.0.0.1:
Traceback (most recent call last):
  File "/usr/local/spark/python/lib/py4j-0.10.6-src.zip/py4j/java_gateway.py", line 852, in _get
    connection = self.deque.pop()
IndexError: pop from an empty deque
```

During handling of the above exception, another exception occurred:

```
Traceback (most recent call last):
  File "/usr/local/spark/python/lib/py4j-0.10.6-src.zip/py4j/java_gateway.py", line 990, in start
    self.socket.connect((self.address, self.port))
ConnectionRefusedError: [Errno 111] Connection refused
```

```
IndexErrorTraceback (most recent call last)

/usr/local/spark/python/lib/py4j-0.10.6-src.zip/py4j/java_gateway.py in _get_connection()
```

```

851         try:
--> 852             connection = self.deque.pop()
853         except IndexError:

```

IndexError: pop from an empty deque

During handling of the above exception, another exception occurred:

```

ConnectionRefusedErrorTraceback (most recent call last)

/usr/local/spark/python/lib/py4j-0.10.6-src.zip/py4j/java_gateway.py in start(self)
989         try:
--> 990             self.socket.connect((self.address, self.port))
991             self.is_connected = True

```

ConnectionRefusedError: [Errno 111] Connection refused

During handling of the above exception, another exception occurred:

```

Py4JNetworkErrorTraceback (most recent call last)

<ipython-input-69-1c9b50703cc7> in <module>()
      5 from pyspark.sql import SparkSession
      6
----> 7 spark = SparkSession.builder.master("local[*]").getOrCreate()

/usr/local/spark/python/pyspark/sql/session.py in getOrCreate(self)
181         session = SparkSession(sc)
182         for key, value in self._options.items():
--> 183             session._jsparkSession.sessionState().conf().setConfString(key,
184             for key, value in self._options.items():
185                 session.sparkContext._conf.set(key, value)

/usr/local/spark/python/lib/py4j-0.10.6-src.zip/py4j/java_gateway.py in __call__(self, *
1156         proto.END_COMMAND_PART
1157
-> 1158         answer = self.gateway_client.send_command(command)
1159         return_value = get_return_value(
1160             answer, self.gateway_client, self.target_id, self.name)

```

```

/usr/local/spark/python/lib/py4j-0.10.6-src.zip/py4j/java_gateway.py in send_command(self)
904         if `binary` is `True`.
905         """
--> 906         connection = self._get_connection()
907         try:
908             response = connection.send_command(command)

/usr/local/spark/python/lib/py4j-0.10.6-src.zip/py4j/java_gateway.py in _get_connection(self)
852         connection = self.dequeue.pop()
853         except IndexError:
--> 854         connection = self._create_connection()
855         return connection
856

/usr/local/spark/python/lib/py4j-0.10.6-src.zip/py4j/java_gateway.py in _create_connection(self)
858         connection = GatewayConnection(
859             self.gateway_parameters, self.gateway_property)
--> 860         connection.start()
861         return connection
862

/usr/local/spark/python/lib/py4j-0.10.6-src.zip/py4j/java_gateway.py in start(self)
995         "server ({0}:{1})".format(self.address, self.port)
996         logger.exception(msg)
--> 997         raise Py4JNetworkError(msg, e)
998
999     def close(self, reset=False):

```

Py4JNetworkError: An error occurred while trying to connect to the Java server (127.0.0.1:4545).

```

In [ ]: # read csv file
data = spark.read.csv('sensor_data.csv',header=True)

In [ ]: from pyspark.sql.types import DoubleType, IntegerType

#convert all columns
for col_name in data.columns:
    data = data.withColumn(col_name, data[col_name].cast(DoubleType()))

In [ ]: # inspect the first 10 rows
data.show(10)

```

```
# the printSchema() method tells you the data type of each column
data.printSchema()
```

1.3.1 Prepare data for model

```
In [ ]: # Split the data into training and test sets (30% held out for testing)
        (trainingData, testData) = data.randomSplit([0.7, 0.3])
```

```
In [ ]: from pyspark.ml import Pipeline
        from pyspark.ml.feature import OneHotEncoder, VectorAssembler
        from pyspark.ml.regression import DecisionTreeRegressor
        from pyspark.ml.evaluation import RegressionEvaluator
```

```
In [ ]: # assemble variables to one feature column
        assembler = VectorAssembler(
            inputCols = ['Device_ID', 'Temp1', 'Temp2', 'Temp3', 'Temp_Ambient'],
            outputCol = "features")
```

```
#      inputCols = ['Date', 'Device_ID', 'Temp1', 'Temp1_ave', 'Temp1_sd',
#                  'Temp2', 'Temp2_ave', 'Temp2_sd',
#                  'Temp3', 'Temp3_ave', 'Temp3_sd',
#                  'Temp_Ambient', 'Temp_Am_ave', 'Temp_Am_sd']
#      outputCol = "features")
```

```
#define the estimator - decision tree
```

```
dt = DecisionTreeRegressor(labelCol="Device_ID", featuresCol="features")
```

```
# Chain indexers and tree in a Pipeline
```

```
pipeline = Pipeline(stages=[assembler, dt])
```

1.3.2 Fit pipeline and transform data

```
In [ ]: #fit the pipeline
```

```
PipelineModel = pipeline.fit(trainingData)
```

```
# transform using the pipeline
```

```
predictions = PipelineModel.transform(testData)
```

```
# evaluate model fit
```

```
predictions.select("prediction", "Device_ID")
```

```
evaluator = RegressionEvaluator(
```

```
    labelCol="Device_ID", predictionCol="prediction", metricName="rmse")
```

```
rmse = evaluator.evaluate(predictions)
```

```
In [ ]: predictions.show()
```

```
In [ ]: ##Root mean square error
```

```
print(rmse)
```


1.3.3 Kmeans clustering

```
In [ ]: from pyspark.ml.clustering import KMeans

        # Trains a k-means model with 4 clusters.
        kmeans = KMeans(featuresCol='features', predictionCol='prediction',k=4)

        #transform data using pipeline
        pipeline = Pipeline(stages=[assembler, kmeans])

        #fir pipeline
        PipelineModel = pipeline.fit(data)

        # transform using the pipeline
        predictions = PipelineModel.transform(data)

In [ ]: #view result
        predictions.show()
```

1.3.4 END

```
In [ ]: spark.stop()
```

1.4 Step 4: Visualization

Time to get creative. Your final task is to build up a set of visualizations that could let an engineer get a quick overview of the current status of the system. Include the current sensor readings for each device and any metrics you think would be important to display. Choose one device and show more detail - a downsampled graph showing the readings over time, perhaps.

You don't need to have your visualizations update in real time - merely show them as they would be presented at a given instant (i.e. feel free to use all the data you stored in the first section).

```
In [8]: import plotly.graph_objs as go
        import plotly.plotly as py

In [9]: import plotly
        plotly.tools.set_credentials_file(username='dgebort18', api_key='wWcelh20cxQaebCFxNBF')
```

1.4.1 Temperature 1's info - Plotly grapf and Autocorelation Plots

```
In [10]: trace1 = go.Scatter(
            y = df['Temp1'],
            mode='lines',
            name = 'Temp 1 actual',
            marker=dict(
                size='16',
                color = np.random.randn(500),
                showscale=True
            )
        )
```

```

)
trace2 = go.Scatter(
    y = df['Temp1_ave'],
    mode='lines',
    name = 'Temp 1 average',
    marker=dict(
        size='16',
        color = np.random.randn(500),
        showscale=True
    )
)
trace3 = go.Scatter(
    y = df['Temp1_sd'],
    mode='lines',
    name = 'Temp 1 standard deviation',
    marker=dict(
        size='16',
        color = np.random.randn(500),
        showscale=True
    )
)
data = [trace1, trace2, trace3]

layout = dict(
    title='Temperature 1 High and Low values',
    xaxis=dict(
        rangeselector=dict(),
        rangeslider=dict(),
        type='date'
    )
)

#layout = dict(title = 'Temperature 1 High and Low values',
#              xaxis = dict(title = 'Time'),
#              yaxis = dict(title = 'Temperature (degrees C)'),
#              )

fig = dict(data=data, layout=layout)
py.iplot(fig, filename='styled-line')

#url_1 = py.plot(data, filename='temp1_info', auto_open=False)
#py.iplot(data, filename='temp1_info')


```

Out[10]: <plotly.tools.PlotlyDisplay object>


Please use the sliding scale to narrow the search / use magnification on a area of interest (down sampling did not work out for this type of plot)

```
In [14]: import matplotlib.pyplot as pyplot
         from pandas import Series
```

```
In [42]: # The code will create the dataset for Temperature 1 Series.
         tst1 = Series(df['Temp1'].values, index=df.Date)
         plt = tst1[1:100].plot(figsize=(7,5))
         pyplot.show(plt)
```




```
In [39]: from pandas.plotting import lag_plot
         lag_plot(tst1)
         pyplot.show()
```



From the scatter plot I can see Temperature 1's data is spread across the diagonal line of the plot - which again indicates a week correlation


```
In [41]: from pandas.plotting import autocorrelation_plot

         autocorrelation_plot(tst1[1:100])
         pyplot.show()
```



Looking at Temperature 1 there is a very weak correlation with a swing of only 0.25 to - 0.25

```
In [32]: from statsmodels.graphics.tsaplots import plot_acf
plot_acf(tst1, lags=60)
pyplot.show()
```



```

In [34]: from statsmodels.tsa.ar_model import AR
         from sklearn.metrics import mean_squared_error

         # split dataset to get only the values
         ts_values = tst1.values

         train, test = ts_values[1:len(ts_values)-16], ts_values[len(ts_values)-16:]

         # train autoregression
         model = AR(train)
         model_fit = model.fit()
         print('Lag: %s' % model_fit.k_ar)
         print('Coefficients: %s' % model_fit.params)

         predictions = model_fit.predict(start=len(train), end=len(train)+len(test)-1, dynamic=False)
         for i in range(len(predictions)):
             print('predicted=%f, expected=%f' % (predictions[i], test[i]))
         error = mean_squared_error(test, predictions)
         print('Test MSE: %.3f' % error)
         # plot results
         pyplot.plot(test)
         pyplot.plot(predictions, color='red')
         pyplot.show()

```

Lag: 25

```

Coefficients: [  5.52448459e+01   2.23449576e-03  -1.27184661e-03  -1.37110952e-02
   2.38284626e-02  -2.21058245e-02  -1.36697268e-02   1.33099454e-02
  -3.30370467e-02  -4.48351345e-02  -1.67587497e-02  -5.38333508e-04
   2.68968361e-02   2.10936620e-02   2.44511578e-02  -3.94294422e-03
  -8.99362771e-03   8.41885866e-03  -1.20510373e-02  -9.33873038e-03
  -7.74145186e-03  -1.43078687e-02  -2.96654252e-02  -1.30514898e-02
  -1.94202971e-02  -2.11975112e-02]

```

```


predicted=43.572488, expected=36.631970
predicted=45.081649, expected=17.254900
predicted=50.409032, expected=78.207766
predicted=50.213273, expected=33.785923
predicted=51.869925, expected=28.099166
predicted=52.264579, expected=48.343542
predicted=48.321485, expected=21.566449
predicted=46.038349, expected=36.915755
predicted=46.779606, expected=26.057700
predicted=46.034037, expected=25.201797
predicted=45.455520, expected=64.241779
predicted=46.232537, expected=76.657513
predicted=46.065451, expected=62.973846

```

```

predicted=44.706429, expected=40.101864
predicted=45.569028, expected=61.580303
predicted=47.248962, expected=96.633182
Test MSE: 524.303

```



Running the AR model, I can see that lag: 25 was chosen and trained. The 16-hour forecast is printed above. The forecast sort of follows the actual values

1.4.2 Temperature 2's info - Plotly graph and Autocorelation Plots

```

In [48]: trace4 = go.Scatter(
            y = df['Temp2'],
            mode='lines',
            name = 'Temp 2 actual',
            marker=dict(
                size='16',
                color = np.random.randn(500),
                showscale=True
            )
        )
        trace5 = go.Scatter(
            y = df['Temp2_ave'],
            mode='lines',
            name = 'Temp 2 average',

```

```

        marker=dict(
            size='16',
            color = np.random.randn(500),
            showscale=True
        )
    )
    trace6 = go.Scatter(
        y = df['Temp2_sd'],
        mode='lines',
        name = 'Temp 2 standard deviation',
        marker=dict(
            size='16',
            color = np.random.randn(500),
            showscale=True
        )
    )

    data = [trace4, trace5, trace6]

    layout = dict(
        title='Temperature 2 High and Low values',
        xaxis=dict(
            rangeselector=dict(),
            rangeslider=dict(),
            type='date'
        )
    )

    fig = dict(data=data, layout=layout)
    py.plot(fig, filename='styled-line')

```

Out[48]: <plotly.tools.PlotlyDisplay object>

1.4.3 Please use the sliding scale to narrow the search / use magnification on a area of interest (down sampling did not work out for this type of plot)

In [60]: *# The code will create the dataset for Temperature 2 Series.*

```

tst2 = Series(df['Temp2'].values, index=df.Date)
plt = tst2[1:100].plot(figsize=(7,5))
pyplot.show(plt)

lag_plot(tst2)
pyplot.show()


autocorrelation_plot(tst2[1:100])
pyplot.show()



plot_acf(tst2, lags=60)


```



```
pyplot.show()
```







This is a similar picture / indication (Temperature sensor 1) of a week correlation of Temperature sensor 2

In [47]: *# split dataset to get only the values*

```
ts_values = tst2.values
```

```
train, test = ts_values[1:len(ts_values)-16], ts_values[len(ts_values)-16:]
```

```
# train autoregression
```

```
model = AR(train)
```

```
model_fit = model.fit()
```

```
print('Lag: %s' % model_fit.k_ar)
```

```
print('Coefficients: %s' % model_fit.params)
```

```
predictions = model_fit.predict(start=len(train), end=len(train)+len(test)-1, dynamic=False)
```

```
for i in range(len(predictions)):
```

```
    print('predicted=%f, expected=%f' % (predictions[i], test[i]))
```

```
error = mean_squared_error(test, predictions)
```

```
print('Test MSE: %.3f' % error)
```

```
# plot results
```

```
pyplot.plot(test)
```

```
pyplot.plot(predictions, color='red')
```


```
pyplot.show()
```

Lag: 25

Coefficients: [7.24233777e+01 2.23449576e-03 -1.27184661e-03 -1.37110952e-02

2.38284626e-02	-2.21058245e-02	-1.36697268e-02	1.33099454e-02
-3.30370467e-02	-4.48351345e-02	-1.67587497e-02	-5.38333508e-04
2.68968361e-02	2.10936620e-02	2.44511578e-02	-3.94294422e-03
-8.99362771e-03	8.41885866e-03	-1.20510373e-02	-9.33873038e-03
-7.74145186e-03	-1.43078687e-02	-2.96654252e-02	-1.30514898e-02
-1.94202971e-02	-2.11975112e-02]		

predicted=57.929737, expected=50.295167
 predicted=59.589814, expected=28.980390
 predicted=65.449935, expected=96.028542
 predicted=65.234601, expected=47.164515
 predicted=67.056917, expected=40.909083
 predicted=67.491037, expected=63.177896
 predicted=63.153633, expected=33.723094
 predicted=60.642184, expected=50.607330
 predicted=61.457567, expected=38.663470
 predicted=60.637441, expected=37.721976
 predicted=60.001072, expected=80.665957
 predicted=60.855790, expected=94.323264
 predicted=60.671996, expected=79.271230
 predicted=59.177072, expected=54.112051
 predicted=60.125931, expected=77.738333
 predicted=61.973858, expected=116.296500
 Test MSE: 634.406



Running the AR model, I can see that lag: 25 was chosen and trained. The 16-hour forecast is printed above. The forecast sort of follows the actual values

1.4.4 Temperature 3's info

```
In [54]: trace7 = go.Scatter(
    y = df['Temp3'],
    mode='lines',
    name = 'Temp 3 actual',
    marker=dict(
        size='16',
        color = np.random.randn(500),
        showscale=True
    )
)
trace8 = go.Scatter(
    y = df['Temp3_ave'],
    mode='lines',
    name = 'Temp 3 average',
    marker=dict(
        size='16',
        color = np.random.randn(500),
        showscale=True
    )
)
trace9 = go.Scatter(
    y = df['Temp3_sd'],
    mode='lines',
    name = 'Temp 3 standard deviation',
    marker=dict(
        size='16',
        color = np.random.randn(500),
        showscale=True
    )
)

data = [trace7, trace8, trace9]

layout = dict(
    title='Temperature 3 High and Low values',
    xaxis=dict(
        rangeselector=dict(),
        rangeslider=dict(),
        type='date'
    )
)

fig = dict(data=data, layout=layout)
```

```
py.iplot(fig, filename='styled-line')
```

Out[54]: <plotly.tools.PlotlyDisplay object>

1.4.5 Please use the sliding scale to narrow the search / use magnification on a area of interest (down sampling did not work out for this type of plot)


In [62]: *# The code will create the dataset for Temperature 3 Series.*



```
tst3 = Series(df['Temp3'].values, index=df.Date)
plt = tst2[1:100].plot(figsize=(7,5))
pyplot.show(plt)
```



```
lag_plot(tst3)
pyplot.show()
```

```
autocorrelation_plot(tst3[1:100])
pyplot.show()
```

```
plot_acf(tst3, lags=60)
pyplot.show()
```







This picture is slightly different but still a very weak correlation. The data from Temperature sensor 3 is grouped in the corners across the diagonal line of the plot

```
In [61]: # split dataset to get only the values
        ts_values = tst3.values

        train, test = ts_values[1:len(ts_values)-16], ts_values[len(ts_values)-16:]

        # train autoregression
        model = AR(train)
        model_fit = model.fit()
        print('Lag: %s' % model_fit.k_ar)
        print('Coefficients: %s' % model_fit.params)


        predictions = model_fit.predict(start=len(train), end=len(train)+len(test)-1, dynamic=False)
        for i in range(len(predictions)):
            print('predicted=%f, expected=%f' % (predictions[i], test[i]))
        error = mean_squared_error(test, predictions)
        print('Test MSE: %.3f' % error)
        # plot results
        pyplot.plot(test)
        pyplot.plot(predictions, color='red')
        pyplot.show()
```

Lag: 25

```
Coefficients: [ 5.59978137e+01 -1.04120543e-02  2.00623249e-02  4.62024193e-02
 1.73631752e-02  2.89807545e-02  1.95073802e-02  3.82435078e-03
-2.55587717e-02 -3.74439006e-02  1.38307500e-02 -3.67027446e-02
-6.31804168e-03  2.01473526e-02 -1.16677391e-02  4.49016735e-02
 3.68074183e-03 -9.88835703e-03 -1.42628190e-02 -4.71354738e-04
-4.88825260e-02  5.74478786e-03 -8.82449307e-04  1.84671549e-04
-1.30737253e-02  3.28170546e-02]

predicted=63.837555, expected=18.315985
predicted=54.176626, expected=8.627450
predicted=62.896705, expected=39.103883
predicted=68.673615, expected=16.892961
predicted=51.456087, expected=214.049583
predicted=72.322973, expected=24.171771
predicted=57.674683, expected=10.783225
predicted=53.076903, expected=18.457877
predicted=57.968395, expected=213.028850
predicted=57.481478, expected=12.600898
predicted=48.548418, expected=32.120889
predicted=61.918293, expected=38.328756
predicted=53.618871, expected=31.486923
predicted=68.689367, expected=20.050932
predicted=55.217424, expected=30.790151
predicted=65.431021, expected=48.316591
```

Test MSE: 4386.013



Running the AR model, I can see that lag: 25 was chosen again and model was trained on this. The 16-hour forecast is printed above. The forecast deviates quite a bit from the actual values peaking at 4 and 8

1.4.6 Ambient Temperature info - Plotly graph and Autocorelation Plots

```
In [66]: trace10 = go.Scatter(  
    y = df['Temp_Ambient'],  
    mode='lines',  
    name = 'Ambient Temp actual',  
    marker=dict(  
        size='16',  
        color = np.random.randn(500),  
        showscale=True  
    )  
)  
trace11 = go.Scatter(  
    y = df['Temp_Am_ave'],  
    mode='lines',  
    name = 'Ambient Temp average',  
    marker=dict(  
        size='16',
```

```

        color = np.random.randn(500),
        showscale=True
    )
)
trace12 = go.Scatter(
    y = df['Temp_Am_sd'],
    mode='lines',
    name = 'Ambient Temp standard deviation',
    marker=dict(
        size='16',
        color = np.random.randn(500),
        showscale=True
    )
)

data = [trace10, trace11, trace12]

layout = dict(
    title='Temperature 3 High and Low values',
    xaxis=dict(
        rangeselector=dict(),
        rangeslider=dict(),
        type='date'
    )
)

fig = dict(data=data, layout=layout)
py.iplot(fig, filename='styled-line')

```

Out[66]: <plotly.tools.PlotlyDisplay object>

1.4.7 Please use the sliding scale to narrow the search / use magnification on a area of interest (down sampling did not work out for this type of plot)

```



In [67]: # The code will create the dataset for Temperature 3 Series.
tstA = Series(df['Temp_Ambient'].values, index=df.Date)
plt = tst2[1:100].plot(figsize=(7,5))
pyplot.show(plt)



lag_plot(tstA)
pyplot.show()

autocorrelation_plot(tstA[1:100])
pyplot.show()

plot_acf(tstA, lags=60)
pyplot.show()

```





These pictures from the Ambient Temperature sensor indicate a high correlation but the correlation coefficient is still only between 0.25 and -0.25 which says this is a weak correlation. The data spread across the diagonal line of the plot indicating a linear trend.

```
In [68]: # split dataset to get only the values
        ts_values = tstA.values

        train, test = ts_values[1:len(ts_values)-16], ts_values[len(ts_values)-16:]


        # train autoregression
        model = AR(train)
        model_fit = model.fit()
        print('Lag: %s' % model_fit.k_ar)
        print('Coefficients: %s' % model_fit.params)

        predictions = model_fit.predict(start=len(train), end=len(train)+len(test)-1, dynamic=False)
        for i in range(len(predictions)):
            print('predicted=%f, expected=%f' % (predictions[i], test[i]))
        error = mean_squared_error(test, predictions)
        print('Test MSE: %.3f' % error)
        # plot results
        pyplot.plot(test)
        pyplot.plot(predictions, color='red')
        pyplot.show()
```

Lag: 25

```
Coefficients: [ 0.10670338  0.04401479  0.05021367  0.00963759  0.05193998 -0.01580651
 0.04833818  0.06457539 -0.00534046 -0.00145579  0.02052801  0.07627349
 0.03171924  0.03096298  0.03677596  0.08002705  0.10041067 -0.0189653
 0.0385952   0.04453086  0.08712103  0.0145506   0.02567537  0.03542656
 0.07829053  0.06807416]

predicted=24.230786, expected=24.110918
predicted=24.210119, expected=23.741410
predicted=24.255270, expected=24.153455
predicted=24.247411, expected=24.215799
predicted=24.251645, expected=23.768086
predicted=24.274848, expected=24.558904
predicted=24.265577, expected=24.271811
predicted=24.267449, expected=24.489534
predicted=24.291308, expected=23.817173
predicted=24.295784, expected=24.228087
predicted=24.293777, expected=24.248050
predicted=24.266753, expected=24.062023
predicted=24.282027, expected=24.171938
predicted=24.269308, expected=24.406753
predicted=24.321363, expected=24.254122
predicted=24.304326, expected=23.611311
Test MSE: 0.087
```



Running the AR model, I can see that lag: 25 was chosen again and model was trained on this. The 16-hour forecast is printed above. The upward trend of the forecast is appropriate to the data.

1.5 Final thoughts:

This was a very challenging Summative assessment but I learnt allot in terms of managing Time series data. This is a great park of my daily work and discussions I have with my customers.

I did not manage to get data to SQL so I instead followed the advice of the moderators to create and use a csv file as my database.

Spark failed me on many occasions (I have two lines of code that i used - the greyed out one keeps requesting a java file link)

Visualizing the data was great: Plotly took some time to get going and autocorrelation porting was informative to myself.